

# Memory Management

`https://www.isical.ac.in/~mandar/courses.html#os`

Indian Statistical Institute

# Outline/Summary I

- Address spaces and address binding
  - compile-time      ■ load-time      ■ run-time
- Memory management: mapping virtual address to physical addresses
  - contiguous allocation and fragmentation
- Paging
  - paging hardware
  - multi-level page tables
  - protection and sharing
- Swapping

- Demand paging
  - page faults
  - page replacement
    - FIFO
    - optimal
    - LRU

## High level code

```
int i;  
int N = 20;  
char prompt[] = "Enter an integer: "  
int A[MAX_SIZE];
```

```
i = N*N + 3*N
```

## High level code

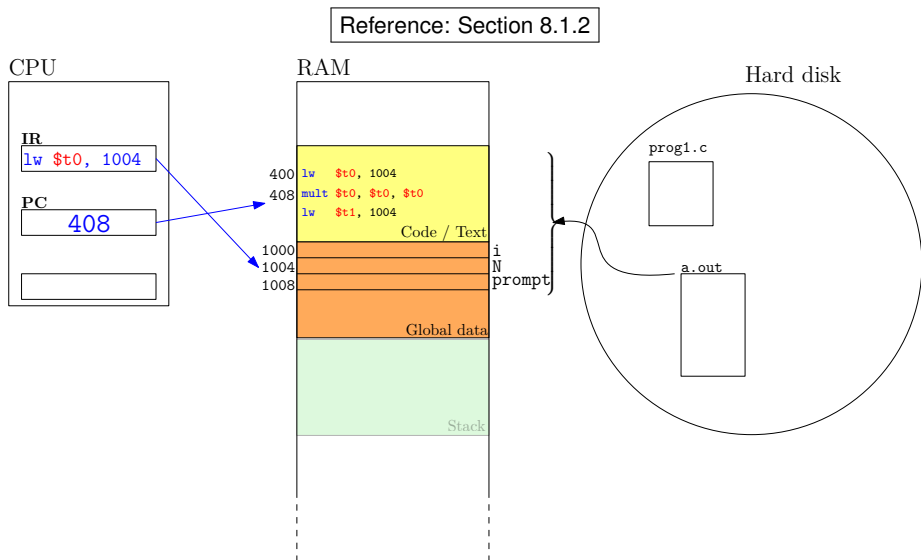
```
int i;  
int N = 20;  
char prompt[] = "Enter an integer: "  
int A[MAX_SIZE];
```

```
i = N*N + 3*N
```

## Equivalent MIPS machine instructions

```
lw    $t0, 1004      # fetch N  
mult  $t0, $t0, $t0   # N*N  
lw    $t1, 1004      # fetch N  
ori   $t2, $zero, 3   # 3  
mult  $t1, $t1, $t2   # 3*N  
add   $t2, $t0, $t1   # N*N + 3*N  
sw    $t2, 1000      # i = ...
```

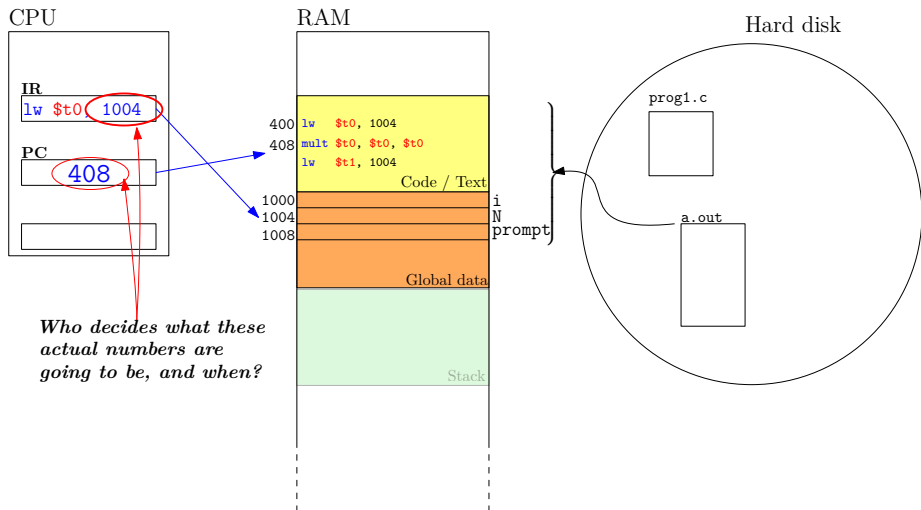
# Address binding



Example adapted from <https://courses.cs.washington.edu/courses/cse378/03wi/lectures/mips-asm-examples.html>

# Address binding

Reference: Section 8.1.2



Example adapted from <https://courses.cs.washington.edu/courses/cse378/03wi/lectures/mips-asm-examples.html>

## Compile-time binding

- Location of program in physical memory must be known at compile time
- Compiler generates *absolute* code
  - compiler binds names to actual physical addresses
- Loading  $\equiv$  copying executable file to appropriate location in memory
- If starting location changes, program will have to be recompiled



## Load-time binding

- Compiler generates *relocatable* code
  - compiler binds names to relative addresses (offsets from starting address)
  - compiler also generates relocation table
- Linker resolves external names and combines object files into one loadable module
- (Linking) loader converts relative addresses to physical addresses
- No relocation allowed during execution

## Run-time binding

- Programs/compiled units may need to be relocated during execution
- CPU generates relative addresses
- Relative addresses bound to physical addresses at runtime based on location of translated units
- Suitable hardware support required


# Memory management unit

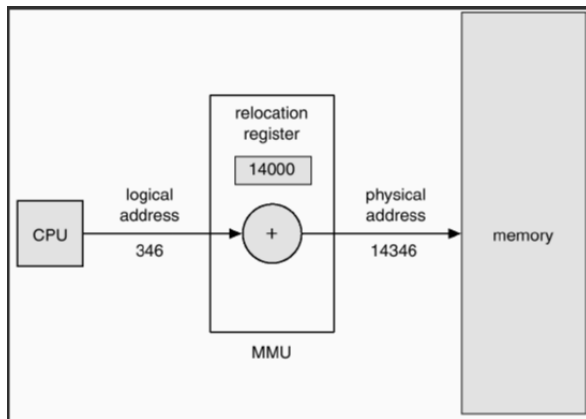
Reference: Section 8.3

- **Logical/virtual address:** address generated by CPU
- **Physical address:** address seen by memory hardware
- Compile-time / load-time binding  $\Rightarrow$  logical address = physical address

Run-time binding  $\Rightarrow$  logical address  $\neq$  physical address

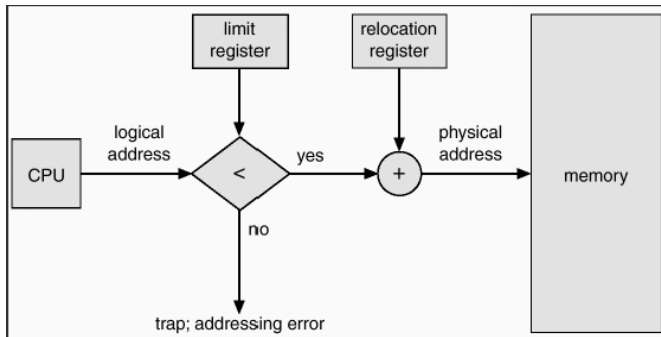
**MMU:** h/w device that maps virtual addresses to physical addresses at run time  
(also called *address translation hardware*)





- Kernel loads relocation register when scheduling a process

# Memory protection



- Prevents process from accessing any memory outside its own address space
- Allows OS size to change dynamically
  - *transient* code (code/data corresponding to infrequently used devices / services) may be removed from memory when not in use

# Contiguous allocation

Reference: Section 8.3

- Memory is divided into variable-sized partitions
- OS maintains a list of allocated / free partitions (*holes*)
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Memory is allocated to processes until requirements of next process in queue cannot be met
  - OS may skip down the queue to allocate memory to a smaller process that fits in available memory
- Hole allocation policies:
  - **First-fit:** allocate the first hole that is big enough
  - **Best-fit:** allocate the smallest hole that is big enough
    - entire free list has to be searched unless sorted
  - **Worst-fit:** allocate the largest hole
- When process exits, memory is returned to the set of holes and merged with adjacent holes, if any

# Contiguous allocation

Example:

Process sizes:


$P_1$  600

$P_2$  1000

$P_3$  300

$P_4$  700

$P_5$  500



# Fragmentation

- **External fragmentation:** memory space to satisfy a request is available, but is not contiguous
  - may be reduced slightly by allocating memory from appropriate end (top/bottom) of hole
- **Internal Fragmentation:** allocated memory may be larger than requested memory
  - ⇒ memory within partition may be left unused
    - may be used to avoid overhead required to keep track of small holes




# Compaction

- Memory contents shuffled to place all free memory together in one large block
- Reduces external fragmentation
- Dynamic relocation (run-time binding) needed

# Paging


Reference: Section 8.5

- Physical memory is partitioned into fixed-size *frames*
- Frame size:
  - defined by hardware
  - should be power of 2
  - typically 512–8192 bytes
- Logical address space is partitioned into *pages* (same size as frames)
- When a process with  $n$  pages has to be loaded,  $n$  free frames have to be found
- Kernel keeps track of free frames
- Page table translates logical page #s to physical frame addresses




# Page table

- Page Table ( $PT$ )  $\equiv$  array of frame addresses
- Address space of process  $P$  has  $k$  pages ( $p_0, p_1, \dots, p_{k-1}$ )  
 $\Rightarrow PT$  for  $P$  has  $k$  (valid) entries
- $PT[i]$  contains starting physical address of frame in which  $p_i$  is stored



# Page table


- Page Table ( $PT$ )  $\equiv$  array of frame addresses
- Address space of process  $P$  has  $k$  pages ( $p_0, p_1, \dots, p_{k-1}$ )  
 $\Rightarrow PT$  for  $P$  has  $k$  (valid) entries
- $PT[i]$  contains starting physical address of frame in which  $p_i$  is stored



## Address translation in MMU



# Address translation



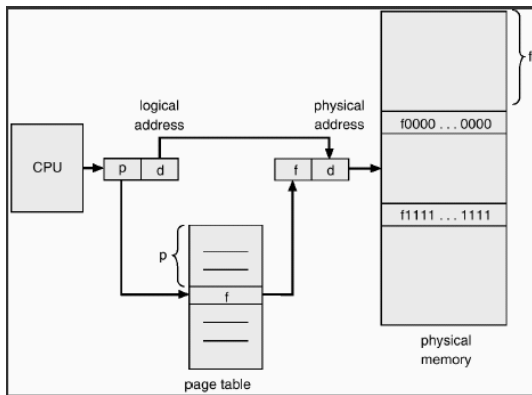
- Which entry of  $PT$  should be consulted?

$$i = \lfloor a/s \rfloor$$

- What byte within the frame pointed to by  $PT[i]$  should be accessed?

$$d = a \bmod s$$

# Address translation



Let  $2^m$  = size of logical address space

$2^n$  = size of page ( $s$ )

Then  $p = m - n$  higher order bits of logical address

$d = n$  lower order bits of logical address

## ■ Page table:

- part of process context
- during context switch, saved page table is used to reconstruct hardware page table
- may be used by some system calls to translate logical addresses to physical addresses in software

## ■ Frame table:

- maintained by kernel
- contains 1 entry per physical page frame
  - whether free or allocated
  - allocation information (PID, page#)

## Miscellaneous issues:

- Memory protection is automatic
  - process cannot address memory outside its own address space
- Fragmentation:
  - No external fragmentation
  - Internal fragmentation can happen
    - half a page per process, on average
- Page/frame size:
  - Small frames  $\Rightarrow$  less fragmentation
  - Large frames  $\Rightarrow$  page table overhead  $\downarrow$ ; I/O is more efficient



## I. **Special purpose registers:**

- Page table is stored in a set of dedicated, high-speed registers
- Instructions to load/modify PT registers are privileged
- Acceptable solution if page table is small
- Example: DEC PDP-11
  - 16-bit address space
  - 8K page size
  - page table contains 8 entries

## II. Memory + PTBR:

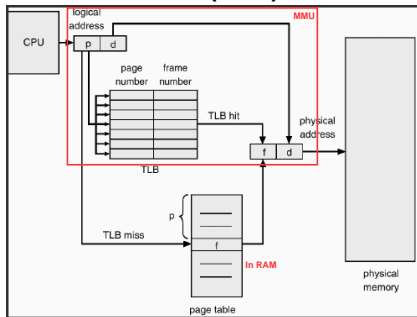
- Needed for large page tables

### Example:

- 32-bit address space
  - page size = 32K =  $2^{15}$  bytes
  - # of entries in page table =  $\frac{2^{32}}{2^{15}} = 2^{17}$
- MMU cannot store so many entries
- PT stored in main memory
  - Base address of PT is stored in *page table base register* (PTBR)  
Length of PT is stored in *page table length register* (PTLR)
  - Context switch involves changing PTBR and PTLR  
(PT stored in memory  $\Rightarrow$  does not need to be saved / restored)
  - Two physical memory accesses are needed per user memory access  
 $\Rightarrow$  memory access is slowed by factor of 2

## III. Associative registers/Translation look-aside buffer (TLB):

- TLB  $\equiv$  small, fast-lookup hardware cache, built using high-speed memory (expensive)
  - each register holds key + value
  - input value is compared *simultaneously* with all keys
  - on match, corresponding value is returned
- TLB holds subset of page table entries
- TLB hit  $\Rightarrow$  additional overhead may be 10% or less  
TLB miss  $\Rightarrow$  new  $\langle$  page#, frame#  $\rangle$  added to TLB
- TLB has to be flushed on context-switch



- **Hit ratio:** percentage of times that a page# is found in TLB
  - depends on size of TLB
- Effective memory access time: average time for a memory access (including TLB lookup)

Example:

TLB lookup: 20ns    Memory access: 100ns    Hit ratio: 80%

Effective access time =  $0.8 \times 120 + 0.2 \times 220 = 140\text{ns}$

# Multi-level paging I

Reference: Section 8.6.1

- Logical address spaces are usually very large ( $2^{32}$  or  $2^{64}$ )
  - ⇒ page tables are very large (how large?)
  - ⇒ page tables should/can not be allocated contiguously

## Example:

- 32-bit address space
- page size =  $4K = 2^{12}$  bytes
- # of entries in page table =  $\frac{2^{32}}{2^{12}} = 2^{20}$
- size of each PT entry = 32 bits =  $2^2$  bytes
  - ⇒ size of PT (in bytes) =  $2^{20} \times 2^2 = 2^{22}$
  - ⇒ size of PT (in pages) =  $2^{22} / 2^{12} = 2^{10}$
  - ⇒ PT is broken into pages and stored non-contiguously

like address space

# Multi-level paging II

- Need a **secondary** PT to keep track of pieces of **primary** PT
- Size of secondary PT =  $2^{10}$  entries =  $2^{12}$  bytes = 1 page

- Two-level paging:


- First level (inner) page table is broken into pieces
- Second level (outer) PT entries point to memory frames holding the pieces of the first level PT

Example: (x86)

← page # →		← offset →
$p_1$	$p_2$	$d$
10 bits	10 bits	12 bits

- 3-, 4-, ... level paging may be required for certain architectures
- Performance:  
for  $k$  levels of paging, TLB miss  $\Rightarrow k$  extra memory accesses

# Memory protection



- Protection bit(s) associated with each frame (via page table entry)
  - protection bit specifies read-only / read-write access
  - protection bit checked in parallel with address computation
  - protection violation (writing to read-only page) causes hardware trap to OS
- Valid/invalid bit indicates whether page is in the process' logical address space
  - set by OS for each page
  - may be used to implement process size restrictions

Reference: Section 8.5.5

- Primarily used for sharing *reentrant* (read-only) code for heavily used programs  
e.g. common utilities, text editors, compilers, window/desktop managers  
NOTE: data for separate processes are stored separately
- PT for each process running a shared program maps code pages to the same physical frames
- Data pages are mapped to different physical frames



Reference: Section 8.3

## Motivation:

Consider the following situation:

$P_1, \dots, P_n$  are resident in memory and occupy all available memory

$P_i$  forks to create a child

## Motivation:

Consider the following situation:

$P_1, \dots, P_n$  are resident in memory and occupy all available memory

$P_i$  forks to create a child

## Principle:


- Space on fast disk (also called **Backing Store**) is used as additional / secondary memory
- Process can be *swapped out* temporarily from main memory to backing store; released memory is used for some other process; swapped process is *swapped in* later for continued execution

# Swapping

## Choosing processes:

### ■ Round-robin

- when  $P$ 's quantum expires, it is swapped out,  $P'$  is swapped into freed memory
- scheduler allocates next quantum to some other process in memory



### ■ Priority-based (**roll out**, **roll in**)

- when higher priority process arrives, lower-priority process is swapped out
- when higher priority process finishes, lower priority process can be swapped in

## Performance:

- Context switch time increases ( $\because$  disk transfer is involved)
- Time quantum should be large compared to swap time for good utilization

## Example:

Process size: 100K    Transfer rate: 1Mbps

$\Rightarrow$  swap-out + swap-in time = 200ms ( $+ \epsilon$ )

## Input/output:

- If  $P$  is swapped out while waiting for input into buffer in user memory, addresses used by I/O devices may be wrong
- Solutions:
  - process with pending I/O should never be swapped, OR
  - I/O operations are always done using OS buffers  
(data can be transferred from OS to user buffer when  $P$  is swapped in)

## Compaction:

1. Processes which have to be moved are swapped out
2. Memory is compacted by merging holes
3. Swapped-out processes are swapped in to different memory locations to minimize fragmentation

Reference: Section 9.1

## Background:

- Instructions being executed /addresses being referenced must be in main memory
- Entire logical address space does not have to be loaded into memory
  - some code may be executed rarely  
e.g. error handling routines for unusual error conditions, code implementing rarely used features
  - arrays/tables may be allocated more memory than required
- Virtual memory  $\equiv$  mechanism to allow execution of processes without requiring the entire process to be in memory

## Advantages:

- Programs can be larger than physical memory
- More programs can be run at the same time  
⇒ throughput / degree of multiprogramming increases without increase in response time
- Less I/O is needed for loading/swapping  
⇒ programs may run faster (compared to swapping)

# Demand paging

Reference: Section 9.2

- Processes reside on secondary memory (high-speed disk)
- When process is to be executed, only the needed pages are brought into memory (**lazy swapping**)
- Page table should specify location of pages (memory vs. on-disk)
  - valid/invalid bit may be used
  - for page that is not currently in memory, page table entry may contain address of page on disk
- While process accesses pages resident in memory, execution proceeds normally
- When process accesses page not in memory, paging hardware traps to OS (**page fault**)



1. Check internal table to determine whether reference was to valid / invalid page.
2. Invalid access  $\Rightarrow$  terminate process.
3. Find a free frame from the free-frame list.
4. Read the desired page from swap device into the free frame.
5. When I/O is complete, update internal table and page table.
6. Restart the instruction that was interrupted by the illegal address trap. (state/context of the process is saved so that process can be restarted in exactly the same state)

## Motivation:

- Pure demand paging: pages are not brought into memory until required  
(process starts executing with no pages in memory)
- Overallocation  $\Rightarrow$  free frame list may be empty when a page fault occurs

## Method:

1. Find the location of the desired page on disk.
2. Find a free frame. If there is no free frame:
  - (a) use page replacement algorithm to select *victim* frame;
  - (b) write victim page to disk; change page/frame tables accordingly.
3. Read the desired page into the (newly) free frame.
4. Update the page and frame tables; restart the process.

# Modify/dirty bit

- Modify/dirty bit is associated with each page (via PT)
- Set whenever the page is written
- If dirty bit of victim frame is clear, it is not written to disk
- Reduces time to service page faults
- Also applicable to read-only pages

# Page replacement algorithms

Reference: Section 9.4

- Page replacement algorithm should yield low page-fault rate
- **Reference string:** sequence of memory references
  - used to evaluate PR algorithms
  - may be generated artificially, or by tracing a process
  - memory references are in terms of page #s only
  - sequence of successive references to the same page may be replaced by only one reference
- # of frames allocated to a process  $\uparrow \Rightarrow$  page faults  $\downarrow$

- Pages are kept in a FIFO queue
  - when a page is brought into memory, it is added at tail of queue
  - when a page has to be replaced, page at head of queue is selected

- Example:

Reference string: 1 2 3 4 1 2 5 1 2 3 4 5

# of frames: 3

# of page faults: 9

- **Belady's anomaly:**

# of frames allocated to a process  $\uparrow \nrightarrow$  page faults  $\downarrow$

**Stack algorithms:**

- Pages in memory with  $n$  frames  $\subseteq$  Pages in memory with  $n + 1$  frames
- Never exhibit Belady's anomaly

# Optimal algorithm

- Replace page that will not be used for the longest period of time
- Minimizes the number of page faults for a fixed number of allocated frames
- Not implementable
- Used to measure other replacement algorithms

# LRU algorithm

- Replace page that has not been used for the longest time
- Often used in practice
- Disadvantage: usually requires substantial hardware assistance

# Examples: FIFO(3)

**Reference string:** 1 2 3 4 1 2 5 1 2 3 4 5

**FIFO** with 3 frames of physical memory allocated

Memory reference	1	2	3	4	1	2	5
Page fault	✓	✓	✓	✓	✓	✓	✓
Victim				1	2	3	4

Pages in memory (first-in)	1	1	1	2	3	4	1
		2	2	3	4	1	2
(last-in)			3	4	1	2	5



## Examples: FIFO(3) (contd.)

### FIFO with 3 frames of physical memory allocated (contd.)

Memory reference	1	2	3	4	5
------------------	---	---	---	---	---

Page fault ✓ ✓

Victim	1	2
--------	---	---

Pages in memory (first-in)	1	1	2	5	5
	2	2	5	3	3
(last-in)	5	5	3	4	4

**No. of PFs: 9**

# Examples: FIFO(4)

**Reference string:** 1 2 3 4 1 2 5 1 2 3 4 5

**FIFO** with 4 frames of physical memory allocated

Memory reference	1	2	3	4	1	2	5
------------------	---	---	---	---	---	---	---

Page fault	✓	✓	✓	✓			✓
------------	---	---	---	---	--	--	---

Victim							1
--------	--	--	--	--	--	--	---

---

Pages in memory (first-in)	1	1	1	1	1	1	2
		2	2	2	2	2	3
			3	3	3	3	4
(last-in)				4	4	4	5

---

# Examples: FIFO(4) (contd.)

**FIFO** with 4 frames of physical memory allocated (contd.)

Memory reference	1	2	3	4	5
Page fault	✓	✓	✓	✓	✓
Victim	2	3	4	5	1

---

Pages in memory (first-in)

3	4	5	1	2
4	5	1	2	3
5	1	2	3	4
1	2	3	4	5

**No. of PFs: 10**

## Examples: LRU(4)

**Reference string:** 1 2 3 4 1 2 5 1 2 3 4 5

## LRU with 4 frames of physical memory allocated

Memory reference    1        2        3        4        1        2        5

Page fault ✓ ✓ ✓ ✓ ✓

Victim 3

## Pages in memory

(LRU)	1	1	1	1	2	3	4
		2	2	2	3	4	1
			3	3	4	1	2
(MRU)				4	1	2	5

# Examples: LRU(4) (contd.)

**FIFO** with 4 frames of physical memory allocated (contd.)

Memory reference    1        2        3        4        5

Page fault                            ✓        ✓        ✓

Victim                                4        5        1

Pages in memory

(LRU)	<div>4</div>	<div>4</div>	<div>5</div>	<div>1</div>	<div>2</div>
	<div>2</div>	<div>5</div>	<div>1</div>	<div>2</div>	<div>3</div>
	<div>5</div>	<div>1</div>	<div>2</div>	<div>3</div>	<div>4</div>
	<div>1</div>	<div>2</div>	<div>3</div>	<div>4</div>	<div>5</div>
(MRU)	<div>1</div>	<div>2</div>	<div>3</div>	<div>4</div>	<div>5</div>

**No. of PFs: 8**

## Stack implementation:

- page numbers are maintained in a doubly-linked stack with *head* and *tail* pointers
- on a page reference, the corresponding PT entry is moved to top of stack
  - six pointers have to be changed
- *tail* points to LRU page

## Background:

- Many architectures do not provide hardware support for true LRU page replacement
- Approximate versions of LRU have to be implemented with the limited hardware support

## Reference bit:

- Associated with each PT entry
- All reference bits are initially cleared by OS
- Set by hardware on each page reference  
⇒ distinguishes used pages from unused pages

## I. **Additional-reference-bits algorithm:**

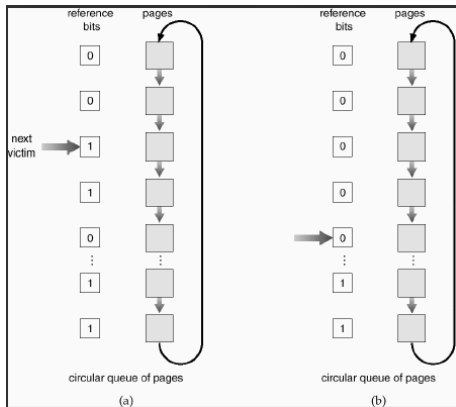
- 1 reference byte associated with each PT entry
- On each timer interrupt: reference byte is right-shifted; reference bit is copied into high-order bit of reference byte and cleared
- Reference bytes contain history of page use for 8 most recent intervals
- Reference bytes order PT entries in LRU order (ties may be broken using FIFO ordering)



# LRU approximation algorithms

## II. Second-chance/clock algorithm:

- Store PT entries in a FIFO queue
- If reference bit of selected page is set:
  - clear reference bit
  - set arrival time to current time
  - continue to next page in FIFO order
- If all bits are set, second-chance replacement reduces to FIFO replacement



## III. Enhanced second-chance algorithm:

- $\langle \text{ref bit}, \text{dirty bit} \rangle$  considered as an ordered pair
  - $\langle 0, 0 \rangle$  – best page to replace
  - $\langle 0, 1 \rangle$  – not recently used, but modified (has to be written to disk)
  - $\langle 1, 0 \rangle$  – recently used, but clean (likely to be used again soon)
  - $\langle 1, 1 \rangle$  – recently used and modified
- First page in lowest non-empty class is selected as victim

# Counting algorithms

- Each PT entry stores count of the number of references to that page
- **LFU Algorithm:** replaces page with smallest count
  - counter may be right shifted at intervals to form an exponentially decaying average usage count
- **MFU Algorithm:** replaces page with largest count
  - LFU page may have been brought in very recently and is yet to be used
- Performance is not very good

# Global vs. local replacement

- Global replacement
  - replacement frame can be selected from all frames (including frames allocated to other processes)
  - generally provides better throughput
- Local replacement: replacement frame can be selected from the frames allocated to the current process

Reference: Section 9.6

## Single user system:

- Kernel occupies  $M$  frames + some frames for dynamic data structures
- Remaining frames are put on free list for use by a user process

## Multiprogramming:

- Minimum # of frames to be allocated to a process:
  - maximum number of memory references permitted in a single instruction

Example: PDP-11 MOV instruction

- instruction may occupy  $> 1$  word
  - 2 operands each of which can be an indirect reference
- if fewer frames are allocated, process should be swapped out, and allocated frames freed

# Allocation of frames

Let  $n$  = # of processes

$M$  = total # of memory frames

$s_i$  = size of process  $p_i$

$a_i$  = # of frames allocated to  $p_i$

**Equal allocation:**

$$a_i = M/n$$

**Proportional allocation:**

$$a_i = M \times s_i / \sum s_i$$

**Priority-based allocation:**


$$a_i = f(P_i, M \times s_i / \sum s_i)$$

**Definition:** situation in which a process is spending more time paging than executing


**Scenario I:**

- Process is not allocated “enough” frames to hold all pages that are in active use
- On a page fault, an active page ( $p$ ) is replaced  
⇒ process page faults soon to page in  $p$

# Thrashing

## Scenario II:

- OS monitors CPU utilization to determine degree of multiprogramming
- Global page replacement algorithm is used
- Process enters a phase where it needs a significantly larger # of frames
- Multiple processes start page-faulting
  - ⇒ paging device queue becomes longer, ready queue empties
  - ⇒ CPU utilization decreases
  - ⇒ CPU scheduler increases degree of multiprogramming





## Local/priority page replacement:

- + If one process starts thrashing, it cannot cause other processes to start thrashing
- Thrashing processes use paging device heavily  
⇒ average service time for page fault increases for non-thrashing processes also

## Page fault frequency monitoring:

- Upper and lower bounds on “desired” page fault rate are determined
- If  $PFR > \text{upper limit}$ , process is allocated another frame  
If  $PFR < \text{lower limit}$ , a frame is removed from the process
- If PFR increases and no free frames are available:
  - a process is selected and suspended
  - freed frames are distributed to processes with high PFRs

## Locality model:

- a set of pages that are actively used together  
e.g. subroutine code, local variables, and some subset of global variables
- process moves from one locality to another (possibly overlapping) locality during execution

## Working set model:

- *Working set window* = most recent  $\Delta$  page references
- *Working set* = set of pages in the working set window
  - approximates the program's current locality
  - $\Delta$  too large  $\Rightarrow$  working set overlaps several localities
  - $\Delta$  too small  $\Rightarrow$  working set does not cover entire locality
- Total demand for frames  $D = \sum WSS_i$

## **Working set model:** (CONTD.)

- OS monitors working set of each process and allocates enough frames to accomodate working set
- If extra frames are available, more processes can be loaded into memory  
If  $D$  exceeds # of available frames, process(es) must be suspended
- Implementation:
  - Timer interrupt is generated at regular intervals e.g. every 5000 memory references
  - For each page, reference bit is copied into history register and cleared
  - Overhead = Frequency of interrupt, # of history bits

Reference: Section 9.3

- Effective access time =  $ma + p \times PF \text{ time}$   
where  $ma$  - memory access time  
 $p$  - probability of a page fault
- Page fault service time:
  - time to service page fault interrupt
  - time for I/O
  - time to restart process

Example:  $PF \text{ Time}$ : 25ms     $ma$ : 100ns

$$EAT \approx 100 + 25,000,000 \times p$$

(for acceptable performance,  $< 1$  memory access in 2,500,000 should fault)

## Swap space:

- Swap space should be allocated in large blocks  
⇒ Disk I/O to swap faster than I/O to file system
- File image can be copied to swap space at process startup
- If swap space is limited: (e.g. BSD UNIX)
  - pages are brought in from file system on demand
  - replaced pages are written to swap space

# Page buffering

- Systems may maintain a pool of free frames
  - On a page fault:
    - required page is read into a free frame from the pool
    - in parallel, a victim is selected and written to disk
    - victim frame is added to free-frame pool
  - Process restarts as soon as possible
  - Page information may also be maintained for each free frame
    - if desired page is in free-frame pool, no I/O is necessary
    - used on VAX/VMS systems with FIFO page replacement
- 
- System may maintain a list of modified pages
  - When paging device is idle, modified pages are written to disk