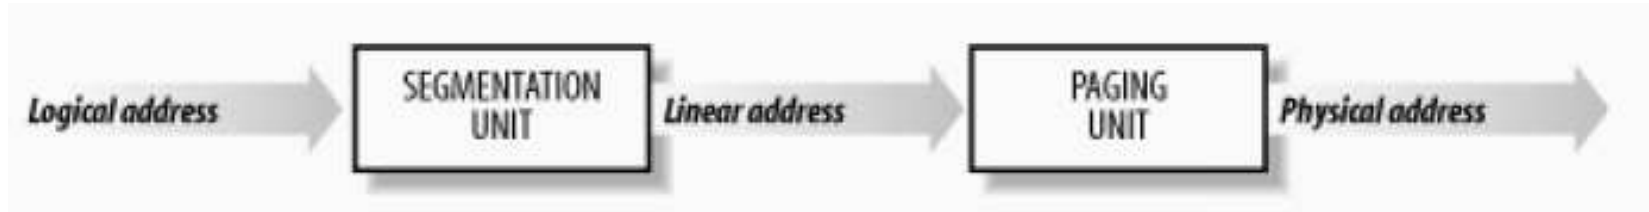


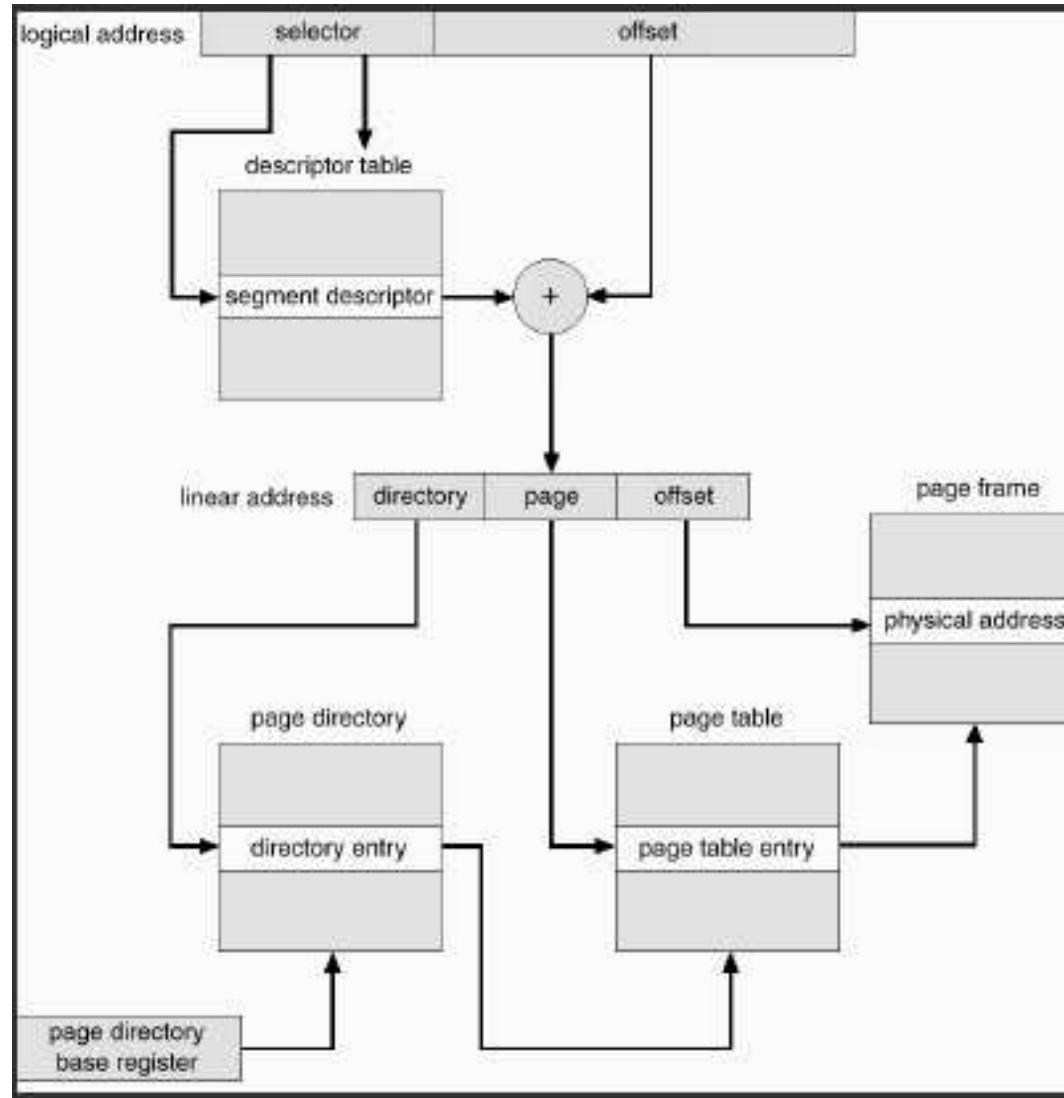
Memory management in i386

Addresses:



- Logical addresses: part of machine language instruction that specifies address of operand / instruction
 - 48-bit 2-tuple ($\langle \text{segment selector}(16), \text{offset}(32) \rangle$)
- Linear addresses: 32-bit unsigned integer
- Physical addresses: bit-pattern put on address lines on the bus

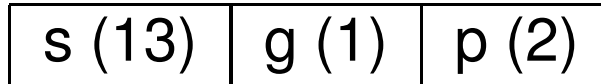
Memory management in i386



Segmentation in i386

- Logical address space is divided into 2 partitions (upto 8K private segments, upto 8K shared segments)
- Segment descriptor tables:
 - segmentation information stored in memory in *Local Descriptor Table* (LDT), *Global Descriptor Table* (GDT)
 - base address of descriptor tables contained in `gdt_r`, `ldt_r` registers

Segment selector



- 13-bit index to identify a segment descriptor entry
- 1-bit table indicator to specify the table to be used (LDT/GDT)
- 2-bit privilege level to hold the current privilege level of CPU

Segment descriptor tables

- *base*(32) – linear address of the first byte of segment
- *limit*(20) – segment length (upto 4 GB)
- *granularity*(1) – whether segment size is expressed in bytes or in 4KB units
- *descriptor privilege level*(2) – minimal CPU privilege level required for accessing the segment
- *segment-present flag*(1) – whether segment is currently in main memory (always 1 in Linux)
- *type*(4) – code, data/stack, etc.

Segmentation registers

- 6 segment registers to hold segment selectors
 - `cs`, `ds`, `ss`: code, data, stack segment
 - `es`, `fs`, `gs`: general purpose (may point to arbitrary data segments)
 - process can address upto 6 segments at one time
- 6 non-programmable 8-byte registers to hold corresponding descriptors from LDT/GDT
 - speeds up translation of logical addresses into linear addresses
 - when a Segment Selector is loaded into segmentation register, corresponding Segment Descriptor is loaded from memory into the matching 8-byte register

Segmentation in Linux

- Linux (+ other variants of UNIX) “ignores” segmentation
(\because RISC architectures have limited support for segmentation)
 \Rightarrow processes see a flat address space
- All segment descriptors are stored in GDT (LDTs not used)

Segment	Base	Limit	DPL	Perm
Kernel code	0x00000000	0xffffffff	0 (kernel mode)	r-x
Kernel data	"	"	"	rW-
User code	"	"	3 (user mode)	r-x
User data	"	"	"	rW-

- Page size: 4K
- 2-level paging is used
- Offset:

p_1 (10)	p_2 (10)	d (12)
------------	------------	----------

 p_1 : Page Directory pointer p_2 : Page Table pointer
- Physical address of page directory is stored in `cr3`
 - `cr3` is changed on a context switch
 - change of `cr3` automatically flushes TLB
- Page table pages:
 - allocated RAM only when required
 - can be swapped to disk

NOTE: Paging is enabled by setting the `PG` flag of `cr0`
(`PG = 0` \Rightarrow linear addresses are interpreted as physical addresses)

Page directory/page table entries:

- *Present flag*: indicates whether page/PT is contained in memory / disk
 - remaining bits can be used to specify location on disk
- *Frame address*: 20 most significant bits of a page frame physical address
- *Reference bit*: set whenever corresponding page frame is addressed*
- *Dirty bit* (only for PT entries): set whenever page frame is written*
- *Permissions*: access rights (read/write or read)
- *User/Supervisor flag*: privilege level required to access the page or Page Table

(*): has to be cleared by OS (not by hardware)

Process address space

- Linear address space of a process is divided into two parts:
 - 3GB (`0x00000000` to `0xbfffffff`)
 - can be addressed in user / kernel mode
 - covered by first 768 entries of PGD
 - 1GB (`0xc0000000` to `0xffffffff`)
 - can be addressed only in kernel mode
 - covered by remaining entries of PGD
 - identical to corresponding entries of the kernel master PGD (same for all processes)
- simplifies user → kernel transitions