

# ***SysV semaphores***

# Data structures

Reference: Stevens(APUE) 14.8

```
struct sem {
    ushort  semval;    // semaphore value
    ushort  semzcnt;   // # waiting for semval == 0
    ushort  semncnt;   // # waiting for increase in semval
    pid_t    sempid;   // process that did last op
};

struct semid_ds {
    ...
    struct sem *sem_base; // array of semaphores (kernel space)
    ushort sem_nsems;      // # of semaphores in the array
    /* timing information */
}
```

# System calls

```
int semid = semget(key_t key, int nsems, int flag);
```

key

- if key is set to IPC\_PRIVATE, new semaphore is created
- if no semaphore corresponding to this key value exists and IPC\_CREAT is asserted in flag, new semaphore is created
- otherwise, integer identifier for the existing semaphore corresponding to key is returned

nsems

- # of semaphores to be created
- 0 if existing semaphore is being used

flag

- specifies permissions, whether semaphore is to be created, etc.

# System calls

```
int semop(int semid, struct sembuf *ops, size_t nops);
```

ops

- array of semaphore operations

```
struct sembuf {  
    ushort sem_num; // which semaphore (0 .. nsems-1)  
    short  sem_op;  
    short  sem_flg; // specifies IPC_NOWAIT, SEM_UNDO, etc.  
}
```

- `sem_op > 0` add this value to `semval` ( $\equiv$  unlocking / returning resources)
- `sem_op < 0` subtract this value from `semval` ( $\equiv$  obtaining resources)
  - `IPC_NOWAIT` (no blocking)  $\Rightarrow$  return -1 (error)
  - otherwise, `semncnt` is incremented and process blocks until `semval >= |sem_op|`
- `sem_op == 0` wait till `semval == 0` (as above)

# System calls

```
int semctl(int semid,int semnum,int cmd,union semun arg);
```

semnum

- specifies which semaphore is to be operated on (0 .. nsems-1)

cmd

- GETVAL gets the value of `sem_base[semnum].semval`
- SETVAL sets the value of `sem_base[semnum].semval` to `arg.val`
- GETALL, SETALL operates on all semaphores using `arg.array`

arg

- specifies values used by various operations

```
union semun {  
    int val;  
    ushort *array;  
    struct semid_ds *buf;  
}
```

# Example application: file locking

Reference: Stevens(UNP) 3.10

- Locking operation: wait till semaphore is free (0); increment by 1
- Unlocking operation: decrement by 1

```
#define SEMKEY 123456L /* arbitrary key value for semget() */

static struct sembuf op_lock[2] = {
    /* semnum, sem_op, sem_flg */
    0,      0,      0,    /* wait for sem#0 to become 0 */
    0,      1,      0    /* then increment sem#0 by 1 */
};

static struct sembuf op_unlock[1] = {
    0, -1, IPC_NOWAIT    /* decrement sem#0 by 1 */
};
```

# Example applications: file locking

```
my_lock()
{
    if (semid < 0) {
        if ((semid = semget(SEMKEY, 1, IPC_CREAT | PERMS)) < 0)
            err_sys("semget error");
    }
    if (semop(semid, &op_lock[0], 2) < 0)
        err_sys("semop lock error");
}

my_unlock()
{
    if (semop(semid, &op_unlock[0], 1) < 0)
        err_sys("semop unlock error"); // "impossible"
}
```

- What happens if a process exits while holding a lock?
  - use `SEM_UNDO` with `semop()`
- What happens when all processes using a semaphore exit?
  - use `IPC_RMID` with `semctl()`
- How are semaphores initialized?
  - use `SETVAL` with `semctl()`
  - race conditions are possible



# ***POSIX semaphores***

```
#include <semaphore.h>
```

```
sem_t mutex;
```

- *Semaphores can only have non-negative values*

- `sem_init(sem_t *sem, int pshared, unsigned int value);`

- `sem_wait(sem_t *sem)` : decrement if semaphore's value is greater than zero, otherwise block until it becomes possible to perform the decrement

- `sem_post` : increments semaphore and another process blocked in `sem_wait()` is woken up and proceeds to lock

- `sem_getvalue`

- `sem_destroy`