

Process Synchronization

Mandar Mitra

Indian Statistical Institute

Cooperating processes

Reference: Section 4.4

Cooperating process: shares data with other processes

Independent process: does not share data with other processes

Means of cooperation:

- Synchronization
- Communication

Producer-consumer problem

Problem: a *producer* process generates output that is used by a *consumer* process as input

Examples:


- print program (producer) + printer driver (consumer)
- `$ ls -l | more`

Implementation:

- Producer and consumer access a shared buffer
- Concurrent access is allowed
- Consumer must wait if buffer is empty
- **Unbounded buffer:** no limit on size of buffer
- **Bounded buffer:** buffer is of fixed size
 - producer must wait if buffer is full


Producer-consumer implementation

- Initially, $in = out = 0$
- Empty queue: $in == out$
full queue: $in+1 \% n == out$
- Buffer can hold at most $n-1$ items



Producer-consumer implementation

- Initially, $in = out = 0$
- Empty queue: $in == out$
full queue: $in+1 \% n == out$
- Buffer can hold at most $n-1$ items



producer

```
p = produce();  
while (in+1 % n == out);  
/* buffer full, skip */  
buffer[in] = p;  
in = in+1 % n;
```

consumer

```
while (in == out); // skip  
p = buffer[out];  
out = out+1 % n;  
consume(p);
```

Race condition

Reference: Section 6.1

producer

```
p = produce();  
while (count==n); // full  
buffer[in] = p;  
in = in+1 % n;  
count++;
```

consumer

```
while (count==0); // empty  
p = buffer[out];  
out = out+1 % n;  
count--;  
consume(p);
```

```
MOV count R0 /* R0 = 5 */  
ADD #1 R0    /* R0 = 6 */  
MOV R0 count /* count=6 */
```

```
MOV count R0 /* R0 = 5 */
```

```
SUB #1 R0    /* R0 = 4 */  
MOV R0 count /* count=4 */
```

Race condition

Reference: Section 6.1

producer

```
p = produce();  
while (count==n); // full  
buffer[in] = p;  
in = in+1 % n;  
count++;
```

consumer

```
while (count==0); // empty  
p = buffer[out];  
out = out+1 % n;  
count--;  
consume(p);
```

```
MOV count R0 /* R0 = 5 */  
ADD #1 R0    /* R0 = 6 */  
MOV R0 count /* count=6 */
```

```
MOV count R0 /* R0 = 5 */
```

```
SUB #1 R0    /* R0 = 4 */  
MOV R0 count /* count=4 */
```

Race condition: several processes manipulate the same data concurrently s.t. outcome depends on the order in which the processes execute

Critical section problem (CSP)

Reference: Section 6.2

Critical section (CS): segment of code in which processes access shared data

Synchronization scheme:

```
while (1) {  
    entry_section();  
    critical_section();  
    exit_section();  
    remainder_section();  
}
```


Desiderata:

- Mutual exclusion: At most one process may execute code from CS at any given time
- Progress: If no process is executing in CS and some processes want to enter CS, only processes not in the remainder section can participate in deciding which process next enters CS. Selection of a process cannot be postponed indefinitely.
- Bounded waiting: there exists a bound on # of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted

Assumptions:

- Basic machine instructions (load, store, test) are executed atomically

Two-process solutions

Algorithm 1: strict alternation

```
shared int turn = 0;
```

```
P_i()  
{ while (1) {  
    while (turn != i); // wait  
    critical_section();  
    turn = j; // j = 1-i  
    remainder_section();  
}  
}
```

- Mutual exclusion – satisfied
- Progress – not satisfied

Two-process solutions

Algorithm 2:

```
shared char want[2] = {0,0};
```

```
P_i()  
{ while (1) {  
    want[i] = 1;  
    while (want[j]);  
    critical_section();  
    want[i] = 0;  
    remainder_section();  
}  
}
```

- Mutual exclusion – satisfied
- Progress – not satisfied

Two-process solutions

Algorithm 3:

```
shared char want[2] = {0,0};
shared int turn = 0;

P_i()
{ while (1) {
    want[i] = 1;
    turn = j;
    while (want[j] && turn!=i);
    critical_section();
    want[i] = 0;
    remainder_section();
  }
}
```

- Mutual exclusion, progress, bounded waiting – satisfied

Two-process solutions

Incorrect solutions:

```
shared char want[2] = {0,0};
shared int turn = 0;
```

```
1. P_i()
2. { while (1) {
3.     turn = j;                want[i] = 1;
4.     want[i] = 1;            while (want[j] && turn!=i);
5.     while (want[j] && turn!=i); OR critical_section();
6.     critical_section();      want[i] = 0;
7.     want[i] = 0;            turn = j;
8.     remainder_section();     remainder_section();
9. }
10.}
```

$$P_0: (3) \rightarrow P_1: (3-5) \rightarrow P_0: (4-5)$$

(Lamport's) bakery algorithm

- Each process wanting to enter CS gets a token number
- Process with lowest token number enters CS
- If two processes have same token number, process with lower PID enters CS

Bakery algorithm

```
shared char choosing[N] = {0, ..., 0};
shared int number[N] = {0, ..., 0};

do {
    choosing[i] = 1; // why??
    number[i] = MAX(number[0], number[1], ..., number[n-1]) + 1;
    choosing[i] = 0; // why??
    for (j = 0; j < n; j++) {
        while (choosing[j]) ;
        while ((number[j] != 0) && (number[j],j < number[i],i)) ;
    }
    critical_section();
    number[i] = 0;
    remainder_section
} while (1);
```

Disabling interrupts

- Critical section executes without preemption
- Adopted by non-preemptive Unix kernels

Atomic test-and-set

- TestAndSet primitive:

```
char TestAndSet(char *flag)
{ char rv = *flag; *flag = 1; return rv; }
```

- CSP solution:

```
shared char lock = 0;
...
while (TestAndSet(&lock));  critical_section();  lock = 0;
...
```


Atomic swap

- Swaps the contents of two words atomically

- CSP solution:

```
shared char lock = 0;
```

```
...
```

```
key = 1;
```

```
do Swap(lock,key) while (key != 0);
```

```
critical_section();
```

```
lock = 0;
```

***n*-process CSP:**

```
shared char lock, waiting[N];
...
waiting[i] = 1;
key = 1;
while (waiting[i] && key) { key = TestAndSet(lock); }
waiting[i] = 0;

critical_section();

j = i+1 % n;
while (j!=i and waiting[j]==0) { j = j+1 % n; }
if (j==i) then lock = 0;
else waiting[j] = 0;
```

Semaphores

Reference: Section 6.4

Definition: a counting semaphore S is an integer variable that can be accessed only through two atomic operations, $wait(S)$ (or $P(S)$) and $signal(S)$ (or $V(S)$)

$wait(S)$: while ($S \leq 0$); $S--$;

$signal(S)$: $S++$;

n -process CSP:

```
shared semaphore mutex = 1;
```

```
P_i()
```

```
{ wait(mutex); critical_section(); signal(mutex); ... }
```

Synchronization problems

- Example:

P_1, P_2 are 2 concurrent processes

Statement S_2 in P_2 should be executed after S_1 in P_1

```
shared semaphore synch = 0;
```

```
P1: ... S1; signal(synch); ...
```

```
P2: ... wait(synch); S2; ...
```

- Careless use may lead to deadlock

```
P1: wait(S1); wait(S2); ... signal(S1); signal(S2);
```

```
P2: wait(S2); wait(S1); ... signal(S2); signal(S1);
```

Busy waiting vs. blocking

Spinlocks/busy waiting

- Processes execute a loop while waiting for entry to critical section
⇒ waste of CPU cycles
- Useful in multiprocessor systems if locks are held for short intervals (\therefore context switch can be avoided)

Blocking

```
typedef struct {int value; struct process *L;} semaphore;
```

wait(S):

```
S.value--;  
if (S.value < 0) {  
    add process to S.L;  
    sleep();  
}
```

signal(S):

```
S.value++;  
if (S.value <= 0) {  
    remove process P from S.L;  
    wakeup(P);  
}
```

Implementation issues

- For bounded waiting, L should be maintained as a FIFO queue
- $wait(S)$, $signal(S)$ must be atomic
 - ⇒ two processes cannot be in $wait/signal$ simultaneously
 - ⇒ $wait / signal$ must be implemented as CS (!!)
 - use hardware solutions (disable interrupt, etc.), OR
 - software solutions
 - busy waiting is limited to the CS in $wait / signal$ only
 - CS in $wait/signal$ is short (~ 10 instructions) and occupied for very short periods of time
(CS in application programs may be long and almost always occupied)

Binary semaphores

Definition: semaphore whose integer value can be either 0 or 1

wait_b(S):

```
if (S.value == 0) {  
    add process to S.L;  
    sleep();  
}  
else S.value = 0;
```

signal_b(S):

```
if (S.value == 0) {  
    P = dequeue(S.L);  
    if (P == NULL) S.value = 1;  
    else wakeup(P);  
}
```

Binary semaphores

Implementing counting semaphores:

```
binary_semaphore S1 = 1, S2 = 0; int C = m;
```

wait(S):

```
wait_b(S3); // why?
wait_b(S1);
C--;
if (C < 0) {
    signal_b(S1);
    wait_b(S2);
}
else signal_b(S1);
signal_b(S3);
```

signal(S):

```
wait_b(S1);
C++;
if (C <= 0) signal_b(S2);
signal_b(S1);
```


Producer-consumer problem

Reference: Section 6.5.1

```
semaphore full = 0, empty = n, mutex = 1;
```

producer

```
while (1) {  
    p = produce();  
    wait(empty);  
    wait(mutex);  
    AddToBuffer(p);  
    signal(mutex);  
    signal(full);  
}
```

consumer

```
while (1) {  
    wait(full);  
    wait(mutex);  
    p = RemoveFromBuffer();  
    signal(mutex);  
    signal(empty);  
    consume(p);  
}
```

Readers-writers problem

Reference: Section 6.5.2

- Shared object is accessed by several concurrent processes
 - **Reader:** processes that only read the shared object
 - **Writer:** processes that update (read + write) the shared object
- Synchronization constraint: two or more readers can access shared data simultaneously; any writer must have exclusive access
- Variants:
 - *first* readers-writers problem: readers have (non-preemptive) priority
 - *second* readers-writers problem: writers have (non-preemptive) priority

Readers-writers problem

```
semaphore mutex = 1, write_access = 1, readcount = 0;
```

reader

```
wait(mutex);  
readcount++;  
if (readcount == 1)  
    wait(write_access);  
signal(mutex);
```

```
read();
```

```
wait(mutex);  
readcount--;  
if (readcount == 0)  
    signal(write_access);  
signal(mutex);
```


writer

```
wait(write_access);  
write();  
signal(write_access);
```

Dining philosophers problem

Reference: [C. A. R. Hoare, 1968](#)

```
philosopher()  
{ while(1) {  
    think();  
    get_chopsticks();  
    eat();  
    release_chopsticks();  
}  
}
```



- Only the nearest chopsticks can be used
- Only free chopsticks can be used
- Chopsticks must be acquired one by one

Dining philosophers problem

```
semaphore chopstick[5] = {1,1,1,1,1};

while (1) {
    think();
    wait(chopstick[i]);
    wait(chopstick[(i+1) % 5]);
    eat();
    signal(chopstick[i]);
    signal(chopstick[(i+1) % 5]);
}
```

Deadlock avoidance:

- Philosophers can pick up chopsticks only if both are available, OR
- Odd philosophers pick up left chopstick first;
even philosophers pick up right chopstick first.

Summary

- Manipulation of shared data \Rightarrow race conditions
- Shared data should be accessed within critical sections
- Solutions to CSP:
 - Algorithm 3 (slide 9), Bakery algorithm
 - assume only atomic loads, stores, tests
 - Hardware solutions
 - assume atomic TestAndSet or Swap instructions
 - Semaphores
 - semaphores use one of the above solutions internally to ensure mutually exclusive access to semaphore variable
 - may be implemented as system calls which block if necessary
- Example synchronization problems: producer-consumer, reader-writer, dining philosopher
 - abstractions of problems that occur in real systems