

Introduction

Reference: Vahalia 3.1

- Many applications involve several independent tasks that can be done concurrently

Examples

1. GUI based applications e.g. web browser

Reference: Vahalia 3.1

- Many applications involve several independent tasks that can be done concurrently

Examples

1. GUI based applications e.g. web browser
2. Standard client-server setup: processing client requests concurrently

Reference: Vahalia 3.1

- Many applications involve several independent tasks that can be done concurrently

Examples

1. GUI based applications e.g. web browser
2. Standard client-server setup: processing client requests concurrently
3. Parallelisable programs that can take advantage of multiprocessor architecture e.g. `make` utility

Reference: Vahalia 3.1

- Many applications involve several independent tasks that can be done concurrently

Examples

1. GUI based applications e.g. web browser
 2. Standard client-server setup: processing client requests concurrently
 3. Parallelisable programs that can take advantage of multiprocessor architecture e.g. `make` utility
- Using multiple processes to achieve concurrency is avoidable:
 - memory load/system overhead increases substantially
 - explicit interprocess communication mechanism must be used

- Solution: use lightweight processes / threads
- Thread / lightweight process \equiv sub-processes within a process
- Threads : process = processes : machine
 - if a thread is blocked, another thread can run
 - timesharing + parallel execution on a multiprocessor

Parallelism vs. concurrency

- **Parallelism:** (physical)

- actual degree of parallel execution achieved
- limited by number of physical processors available

- **Concurrency:** (conceptual)


- maximum parallelism achievable with unlimited processors
- determined by application and its design

Parallelism vs. concurrency

		Parallelism →	
		Uniprocessor	Multiprocessor
Concurrency ↓	Single thread	<ul style="list-style-type: none">* Time sharing* Several processes with nearly identical address space	<ul style="list-style-type: none">* Separate processes run in parallel
	Multi thread	<ul style="list-style-type: none">* If one thread blocks, process can continue* Lower memory overhead, easier comm./synch.	<ul style="list-style-type: none">* True parallelism (N threads multiplexed on n processors)* Synchronization harder

Basic features

- Process = set of threads + collection of shared resources
- Shared resources:
 - address space (code + data)
 - user credentials
 - open files
 - child processes
- Private resources for each thread:
 - PC, stack, register context
 - child threads
 - state
- No protection between threads \Rightarrow programmer is responsible for synchronization to prevent data corruption



User threads, kernel threads, LWPs

User threads

Reference: Vahalia 3.2.3

- Threads abstraction provided by user level library
- Library provides functions for creation, destruction, switching, scheduling of threads without kernel support
- Each user thread has:
 - user stack
 - area to save user-level register context
 - signal mask
 - state information, etc.

Can be saved and restored without kernel intervention.

Synchronisation

- Global data structures shared
⇒ must be protected using synchronisation primitives e.g., lock variables/semaphore
- Thread library provides implementation of semaphores (or similar)
- Synchronisation operations can block threads and switch to other thread (if necessary)

User threads: asynchronous I/O

- Allows processes to perform I/O without blocking
- Read/write request simply queues the operation and returns; when I/O completes, process is informed via `SIGPOL`
- Programming using AIO is complex
- Threads library uses asynchronous methods internally and provides applications a synchronous programming environment
 - each request is synchronous w.r.t. calling thread (thread blocks until I/O completes)
 - library invokes asynchronous I/O operation and schedules another thread
 - on I/O completion, library reschedules blocked thread

User threads: advantages

- Natural, synchronous programming model
- Thread operations / interactions involve only user-level context (no system call / kernel mode switch required)
⇒ extremely lightweight (fast, low memory overhead)

Example:

SPARC 2: user thread creation - $50\text{-}60\mu s$
process creation - $1700\mu s$

User threads: disadvantages

- Kernel schedules processes without knowledge of constituent threads / thread-level priorities
 - when process is pre-empted, all its threads are pre-empted
 - process running a high priority user thread may be pre-empted in favour of a process running a low priority thread
- No parallelism even on a multiprocessor
- Thread switching
 - clock interrupts occur periodically \Rightarrow scheduler can be run
 - once a thread starts running, no other thread will run until the thread voluntarily gives up CPU (calls thread library function)

Reference: Vahalia 3.2.1

- Created/destroyed as needed by the kernel for executing a specific function
- *Not visible to user programs*
- Shares kernel text, global data
- Private resources:
 - thread table entry
 - kernel stack
 - register context
 - scheduling / synchronization info
- Relatively inexpensive to create, context switching is quick
(memory mappings do not have to be changed)

Examples: system processes

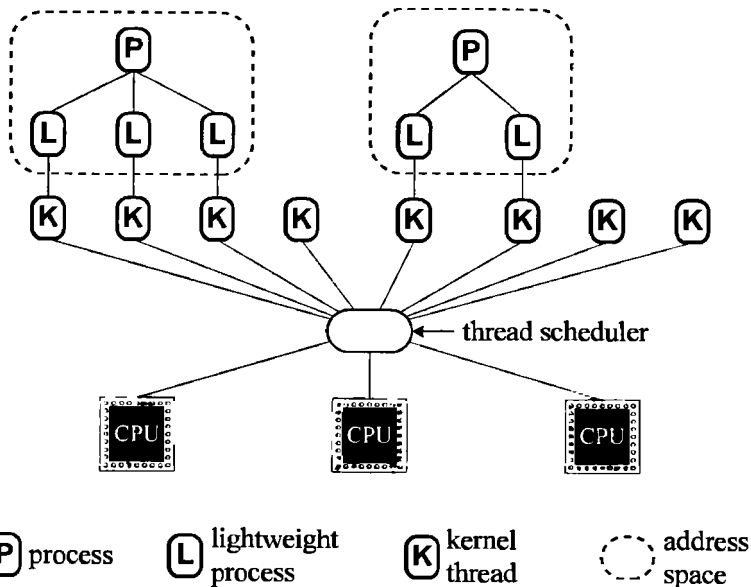
- Implemented as *processes* in traditional systems since there is no provision for kernel threads
 - daemon processes start at user level, but execute entirely in kernel mode \Rightarrow functionally equivalent to kernel threads (process specific admin. info == unnecessary overhead)
- Implemented as *kernel threads* in modern multi-threaded kernels

Lightweight processes

Reference: Vahalia 3.2.2

- LWP \equiv kernel-supported user thread / *kernel thread “visible” to users*
- Process contains one or more LWPs
 - each LWP is supported by a separate kernel thread
 - LWPs share address space and other resources of process

Lightweight processes



Lightweight processes

- LWPs are independently scheduled by kernel
- On a multiprocessor, each LWP can be despatched to run on a different processor
- Resource or I/O wait blocks individual LWPs (not entire process)
- Access to shared data has to be synchronized
 - if an LWP tries to access locked data, it will block / busy-wait
 - busy waiting
 - user mode operation \Rightarrow low overhead
 - good option for small critical sections / resources that are held only briefly

LWP: disadvantages

- Creation/destruction/synchronisation/scheduling of LWPs require system calls
 - mode switch + copy between user and kernel address space required
 - unsuitable for applications that
 - use a large number of threads
 - create/destroy threads frequently
 - control is frequently transferred from thread to thread
- LWPs are useful if each thread is fairly independent of the others
(frequent access of shared data \Rightarrow synchronization overhead \uparrow)

Design issues

Design issues: stack growth

Reference: Vahalia 3.3.5

■ Single-threaded process:

- dedicated stack segment
- stack overflow → protection fault → kernel automatically extends stack (instead of sending a signal)

■ Multi-threaded process:

- several user stacks
- stacks allocated by threads library, possibly from heap/data region
- library may protect against overflow by allocating a write-protected page just beyond the end of stack
- stack overflow → SIGSEGV → thread handles it appropriately

Design issues: fork


Reference: Vahalia 3.3.1

Duplicate all LWPs of parent or only the one that invokes fork?

Option1: Copy only the calling LWP into the new process

Advantages:

- more efficient
- preferable if child calls *exec* to invoke another program after *fork*



Disadvantages: LWPs may be used to support user-level thread libraries (user thread \equiv data structure in user space)

- new process may contain user-level threads that are not bound to any LWP
- if child process tries to acquire locks held by a non-existent thread, deadlock may occur

Design issues: fork

Duplicate all LWPs of parent or only the one that invokes fork?

Option2: Copy entire process (all LWPs)

Advantages:

- preferable when the entire process needs to be cloned (rather than `exec`)

Disadvantages:

- if cloned LWP is manipulating shared data structures, then shared data may become corrupted

Reference: Vahalia 3.3.4


Visibility: Should LWPs be visible outside the process?

- Not visible to other processes
- LWPs within a process can see / signal each other

Solaris threads

Solaris/SVR4 threads

Reference: Vahalia 3.6



Kernel threads:

- Solaris kernel is organized as a set of kernel threads
- Kernel threads independently scheduled / dispatched
- May run LWP or execute internal kernel function (i.e. not associated with any process)
- Fully pre-emptible
- Context switching between threads is less expensive than context switching between processes (virtual address space does not have to be remapped)

Kernel thread specific resources:

- stack, pointer to stack
- saved copy of kernel registers
- priority / scheduling information
- pointers to connect thread record in a scheduler queue / blocked queue
- pointers to associated *lwp* and *proc* structures
- pointers to maintain a queue of all threads in a process, all threads in the system

Lightweight processes:

- Bound to its own kernel thread throughout its lifetime
 - LWPs are scheduled independently and may execute in parallel on multiprocessors
- Traditional *proc* structure + *u area* replaced by:
 - *proc* structure – holds all per-process data (including process specific part of traditional *u area*)
 - *lwp* structure to hold all per-LWP data
 - saved values of user-level registers
 - system call arguments, results, error code
 - signal handling information
 - resource usage, timing information, profiling data
 - alarms
 - pointers to kernel thread structure + parent *proc* structure

User threads:

- Implemented by the threads library
 - provides commonly used API
 - threads created, destroyed, managed without kernel interference
- Run on top of LWPs
 - details taken care of by threads library
 - library creates a pool of LWPs
 - all user threads are multiplexed on this pool of LWPs
 - threads may be *bound* to a dedicated LWP, or *unbound*
 - relation between LWPs and user threads similar to relation between standard I/O library routines (high-level API) and UNIX systems calls (low-level API, more control)

- Single LWP created by kernel when program is started; executes thread compiled as the main program
 - additional threads created by library calls

```
thread_id_t thread_create(char *stack_addr,  
                          u_int stack_size,  
                          void (*func)(), void *arg,  
                          int flag)
```

where `flag` determines whether
new LWP is to be created,
thread is to be permanently bound to this LWP

Thread data structure:

- thread id (allows threads within a process to communicate with each other)
 - saved register state
 - user stack - allocated by library
 - signal mask
 - within process priority - used by thread scheduler (not known to kernel)
 - thread local storage - statically allocated data that is not shared between threads
- ```
#pragma unshared errno
extern int errno;
```

## System calls:

### ■ *fork*

- duplicates each LWP of parent in child
- any LWPs that were in the middle of a system call return with `EINTR` error

### ■ *fork1*

- duplicates only the thread that invokes the function
- useful when child process expects to invoke new program

### ■ *pread, pwrite*

- enables concurrent random I/O by taking seek offset as an argument

### ■ *exec*

- first forces all but the calling LWP to exit